


# End of Moore's Law

## End of Moore's Law

Perhaps the best source to describe what is happening with chips is [Dr. David Patterson](#) a professor at UC Berkeley and an architect for the TPU (Tensorflow Processing Unit).

## End of Growth of Single Program Speed?

40 years of Processor Performance



Based on SPECintCPU. Source: John Hennessy and David Patterson, Computer Architecture: A Quantitative Approach, 6/e. 2018

15

The high-level version is that CPU clock speed is effectively dead. This opens up an opportunity for new solutions. These solutions involve both cloud computing and also specialized chips called [ASICs](#).

## ASICs: GPUs, TPUs, FPGA

### ASIC vs CPU vs GPU

## TPU in Production

# TPU v1: in production since 2015



### Search

Search ranking  
Speech recognition



### Translate

Text, graphic and  
speech translation




### Photos

Photos search




## CPU

### How CPU works



SOFTWARE




OUTPUT



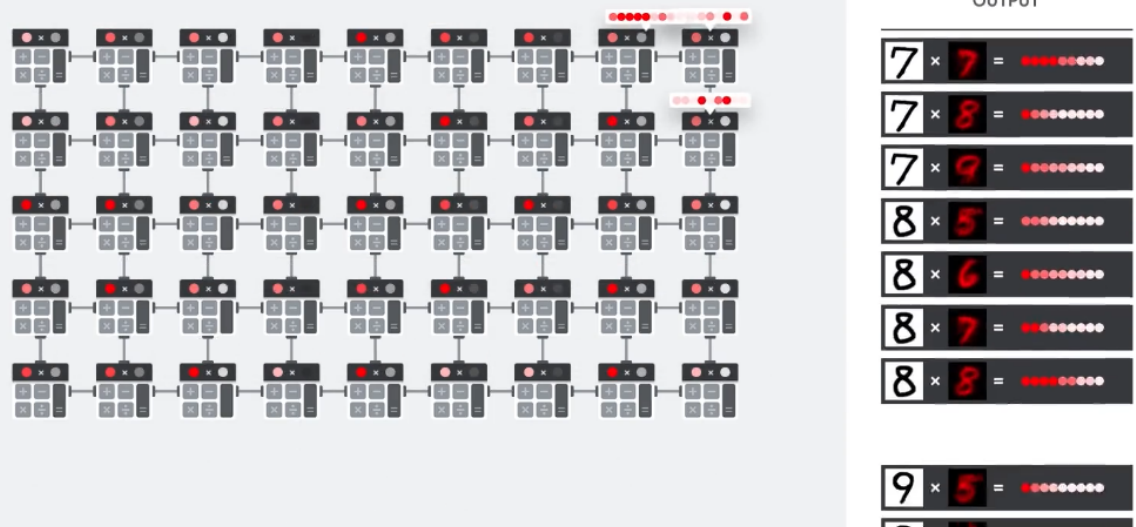
## GPU

## How GPU works



## TPU

## How TPU works




TPU v2 has 32,768 multipliers and adders

Sources: <https://storage.googleapis.com/nexttpu/index.html>

## Using GPUs and JIT

Screencast



One of the easiest ways to use a Just in Time compiler (JIT) or a GPU is to use a library like [numba](#) and a hosted runtime like [Google Colab](#).

There is a step by step example of how to use these operations in the following notebook ([https://github.com/noahgift/cloud-data-analysis-at-scale/blob/master/GPU\\_Programming.ipynb](https://github.com/noahgift/cloud-data-analysis-at-scale/blob/master/GPU_Programming.ipynb)) [[https://github.com/noahgift/cloud-data-analysis-at-scale/blob/master/GPU\\_Programming.ipynb](https://github.com/noahgift/cloud-data-analysis-at-scale/blob/master/GPU_Programming.ipynb)]. The main high level takeaway is that the GPU runtime in colab must be enabled.

# Notebook settings

Runtime type

Python 3

Hardware accelerator

GPU



Runtime shape

Standard

☐ Omit code cell output when saving this notebook

CANCEL

SAVE

Next up install `numba` and double check the CUDA `.so` libraries are available.

```
!pip install numba
!find / -iname 'libdevice'
!find / -iname 'libnvvm.so'
```

You should see something like this.

```
/usr/local/cuda-10.0/nvvm/libdevice
/usr/local/cuda-10.1/nvvm/libdevice
/usr/local/cuda-10.0/nvvm/lib64/libnvvm.so
/usr/local/cuda-10.1/nvvm/lib64/libnvvm.so
```

Next up try one of the methods for speeding up your code.

## GPU Workflow

[TO DO: Create GPU vectorize workflow diagram]

# Exercise

- Topic: Go through [colab example here](#)
- Estimated time: 20-30 minutes
- People: Individual or Final Project Team
- Slack Channel: #noisy-exercise-chatter
- Directions:
  - Part A: Get code working in colab
  - Part B: Make your own GPU or JIT code to speed up a project you are working on. Share in slack and/or create a technical blog post about it.